

SHIPSAFE AI

QUICK START

Install the guardrails and establish launch truth in sixty minutes.

CONTROL THE AGENT. PROTECT THE PROJECT. SHIP WITH EVIDENCE.

Version 1.0 | June 2026

Use this bundle in order

1. Copy the matching agent instruction file and UNIVERSAL_PROJECT_GUARDRAILS.md into the project root.
2. Complete PROJECT_CONTEXT.md, SCOPE_LOCK.md, ACCEPTANCE_CRITERIA.md, and PROTECTED_FILES.md.
3. Create a known-good checkpoint before material work.
4. Open the Launch Readiness Scorecard and replace Unknown only when evidence exists.
5. Fix Critical failures and unknowns before public launch.
6. Use GO_NO_GO_DECISION.md to record the release decision and residual risk.

Minimum project file set

File	Purpose
AGENTS.md / CLAUDE.md / GEMINI.md / Cursor rule	Persistent tool-specific instructions
UNIVERSAL_PROJECT_GUARDRAILS.md	Shared authority, safety, evidence, and stop rules
PROJECT_CONTEXT.md	Architecture, data, commands, and constraints
SCOPE_LOCK.md	What may and may not change
ACCEPTANCE_CRITERIA.md	Observable definition of success
PROTECTED_FILES.md	Known-good behavior and sensitive boundaries
ROLLBACK_PLAN.md	Recovery trigger and procedure

Five things that block launch

- Any Critical check is failed or unknown.
- One user may access another user’s data or files.
- Secrets may be present in source, history, logs, prompts, or browser bundles.
- Failed integrations can appear successful through mocks, static data, or silent fallback.
- There is no tested path back to a known-good state.

Paste this before a risky coding task

Inspect before editing. State the smallest viable plan, affected files, protected behavior, risk, tests, and rollback. Do not invent success, weaken tests/security, expose secrets, or make destructive/production/auth/database/billing/infrastructure changes without explicit approval. Stop when requirements conflict or evidence is missing. Completion requires changed files, commands, test output, acceptance-criteria evidence, remaining risk, and rollback instructions.

Launch decision

Score	Use
95-100%	Launch ready if no Critical failure or unknown remains.
85-94%	Conditional launch with written residual-risk acceptance.
70-84%	Controlled staging only.
Below 70%	Do not launch publicly.

Important

A numerical score never compensates for a Critical failure. SHIPSAFE is a control framework, not a guarantee or security certification.

Fast launch-risk triage

- Cross-user data access has been tested directly.
- Anonymous and privileged API behavior is verified.
- Source, history, logs, prompts, and browser bundles were checked for secrets.
- Mock/static/fallback behavior is inventoried.
- Dependency failure cannot appear successful.
- Critical existing behavior has regression tests.
- A known-good version and tested rollback exist.
- Database and storage access are tested with normal user credentials.
- Agent permissions, protected files, and approval gates are explicit.
- Acceptance criteria have reviewable evidence.

When to stop

Stop public launch when any Critical check is failed or unknown, when data/authorization boundaries are unverified, when secrets may be exposed, when failed dependencies can look successful, or when recovery has not been tested.